

# Rust Trainings All in One

- High-level intro about Rust
- Ownership, borrow check, and lifetime
- Typesystem and data structures
- Concurrency - primitives
- Concurrency - async/await
- Networking and security
- FFI with C/Elixir/Swift/Java
- WASM/WASI
- Rust for real-world problems

# High-level Intro About Rust

# Why Rust?

# Let's talk about values and tradeoffs first

- Approachability
- Availability
- Compatibility
- Composability
- Debuggability
- Expressiveness
- Extensibility
- Interoperability
- Integrity
- Maintainability
- Measurability
- Operability
- Performance
- Portability
- Productivity
- Resiliency
- Rigor
- Safety
- Security
- Simplicity
- Stability
- Thoroughness
- Transparent
- Velocity

# C

- Approachability
- Availability
- Compatibility
- Composability
- Debuggability
- Expressiveness
- Extensibility
- Interoperability

- Integrity
- Maintainability
- Measurability
- Operability
- **Performance**
- Portability
- Productivity
- Resiliency

- Rigor
- Safety
- Security
- **Simplicity**
- Stability
- Thoroughness
- **Transparent**
- Velocity

# Erlang/Elixir

- Approachability
- Availability
- Compatibility
- Composability
- Debuggability
- Expressiveness
- Extensibility
- Interoperability
- Integrity
- Maintainability
- Measurability
- Operability
- Performance
- Portability
- **Productivity**
- **Resiliency**
- Rigor
- **Safety**
- Security
- **Simplicity**
- Stability
- Thoroughness
- Transparent
- Velocity

# Python

- **Approachability**
- Availability
- Compatibility
- Composability
- Debuggability
- **Expressiveness**
- Extensibility
- Interoperability
- Integrity
- Maintainability
- Measurability
- Operability
- Performance
- Portability
- **Productivity**
- Resiliency
- Rigor
- Safety
- Security
- **Simplicity**
- Stability
- Thoroughness
- Transparent
- Velocity

# Java (in early days)

- Approachability
- Availability
- Compatibility
- Composability
- Debuggability
- Expressiveness
- Extensibility
- Interoperability
- Integrity
- Maintainability
- Measurability
- Operability
- **Performance**
- **Portability**
- Productivity
- Resiliency
- Rigor
- **Safety (memory)**
- **Security**
- Simplicity
- Stability
- Thoroughness
- Transparent
- Velocity

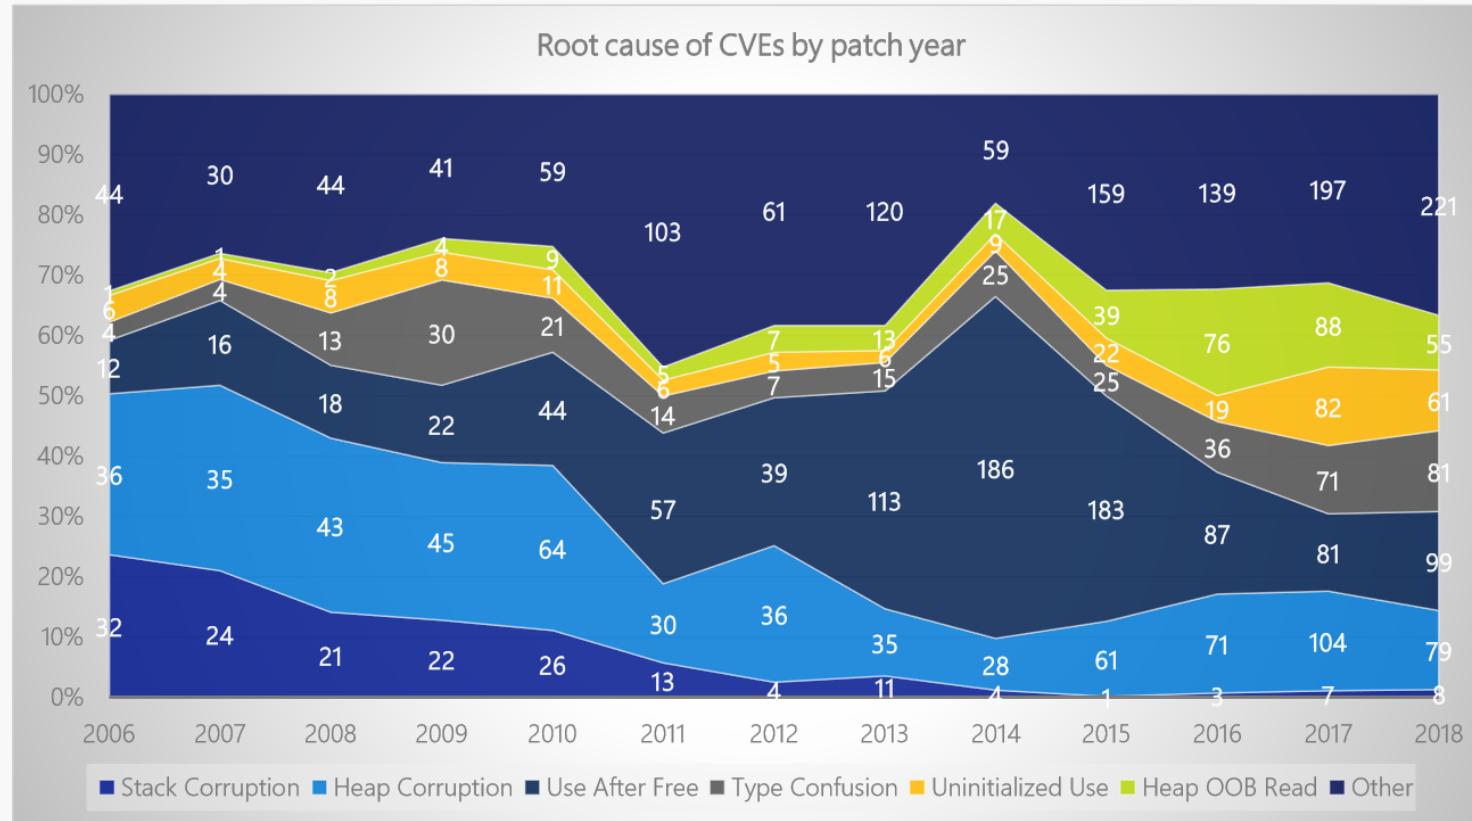
# Rust

- Approachability
- Availability
- Compatibility
- Composability
- Debuggability
- **Expressiveness**
- Extensibility
- Interoperability
- Integrity
- Maintainability
- Measurability
- Operability
- **Performance**
- Portability
- **Productivity**
- Resiliency
- Rigor
- **Safety!!!**
- Security
- Simplicity
- Stability
- Thoroughness
- Transparent
- Velocity



# **Why safety is important?**

# Drilling down into root causes



Stack corruptions are essentially dead

Use after free spiked in 2013-2015 due to web browser UAF, but was mitigated by Mem GC

Heap out-of-bounds read, type confusion, & uninitialized use have generally increased

Spatial safety remains the most common vulnerability category (heap out-of-bounds read/write)

Top root causes since 2016:

#1: heap out-of-bounds

#2: use after free

#3: type confusion

#4: uninitialized use

Note: CVEs may have multiple root causes, so they can be counted in multiple categories

# Safety is hard!

- memory safety is not easy (you need to understand the corner cases)
- concurrency safety is really hard (without certain tradeoffs)
- Often you have to bear the extra layer of abstractions
  - normally it means performance hit

# Memory safety

- Manually - C/C++: painful and error-prone
- Smart Pointers - C++/ObjC/Swift: be aware of cyclical references
- GC - Java/DotNet/Erlang: much bigger memory consumption, and STW
- Ownership - Rust: learning curve

# Concurrency safety

- single-threaded - Javascript: cannot leverage multicore
- GIL - Python/Ruby: multithreading is notorious inefficient
- Actor model - Erlang/Akka: at the cost of memory copy and heap allocation
- CSP - Golang: at the cost of memory copy and heap allocation
- Ownership + Type System - Rust: super **elegant** and **no extra cost!**

**How Rust achieves  
memory and concurrency safety  
without extra cost?**



**Show me the code!**

```
fn main() {
    let mut arr: Vec<i32> = vec![1, 2, 3];      move occurs because `arr` has type `Vec<i32>`, which does not implement the `Copy` trait
    arr.push(4);

    let _result: Result<(), Error> = process(arr);    value moved here
    let _v: Option<i32> = arr.pop(); // failed since arr is moved    borrow of moved value: `arr`

    // you can have multiple immutable references
    let mut arr1: Vec<i32> = vec![1, 2, 3];
    let ir1: &Vec<i32> = &arr1;
    let ir2: &Vec<i32> = &arr1;    immutable borrow occurs here

    println!("ir1: {:?} ir2: {:?}", ir1, ir2);

    // but you can't have both mutable and immutable references
    let mr1: &mut Vec<i32> = &mut arr1;    cannot borrow `arr1` as mutable because it is also borrowed as immutable
    // let mr2 = &mut arr1;

    println!("mr1: {:?} mr2: {:?}", mr1, ir2);    immutable borrow later used here

    // by default, closure borrows the data
    let mut arr2: Vec<i32> = vec![1, 2, 3];
    thread::spawn(|| {    closure may outlive the current function, but it borrows `arr2`, which is owned by the current function
        ... arr2.push(4);    `arr2` is borrowed here
    });
}

// we shall move the data explicitly
let mut arr3: Vec<i32> = vec![1, 2, 3];
thread::spawn(move || arr3.push(4));
}

fn thread_safety() {
    // but certain types cannot be moved to other thread safely
    let mut rc1: Rc<Vec<i32>> = Rc::new(vec![1, 2, 3]);
    thread::spawn(move || {    `Rc<Vec<i32>>` cannot be sent between threads safely
        ... rc1.push(4);
    });
}
```

```
fn thread_safety_reasoning() {
    let mut map: HashMap<&str, &str> = HashMap::new();      move occurs because `map` has type `HashMap<&str, &str>`, which does not implement
    map.insert(k: "hello", v: "world");

    // Arc is an atomic reference counter which can be moved safely across threads
    let mut ir: Arc<HashMap<&str, &str>> = Arc::new(data: map);      variable does not need to be mutable
    map.insert(k: "hello1", v: "world1"); // you can't do this since map is moved      borrow of moved value: `map`
    let ir1: Arc<HashMap<&str, &str>> = ir.clone(); // this is cheap, just reference counter clone
    thread::spawn(move || assert_eq!(ir1.get("hello"), Some(&"world")));
    // but arc is immutable, so this would fail
    thread::spawn(move || ir.insert(k: "hello2", v: "world2"));      cannot borrow data in an `Arc` as mutable

    // the compiler guide you to use types that provides mutable reference for threads

    // use Mutex - you can't clone a Mutex, thus you can't make it available for multiple threads
    let mut map1: HashMap<&str, &str> = HashMap::new();
    map1.insert(k: "hello", v: "world");
    let mr: Mutex<HashMap<&str, &str>> = Mutex::new(map1);
    let mr1 = mr.clone();      no method named `clone` found for struct `Mutex<HashMap<&str, &str>>` in the current scope
    thread::spawn(move || mr.lock().unwrap().insert(k: "hello1", v: "world1"));
    mr1.lock().unwrap().insert("hello2", "world2");

    // use Mutex with Arc - now you have mutable access and multi-thread cloning
    let mut map2: HashMap<&str, &str> = HashMap::new();
    map2.insert(k: "hello", v: "world");
    let mr: Arc<Mutex<HashMap<&str, &str>>> = Arc::new(data: Mutex::new(map2));
    let mr1: Arc<Mutex<HashMap<&str, &str>>> = mr.clone();

    thread::spawn(move || mr.lock().unwrap().insert(k: "hello1", v: "world1"));
    thread::spawn(move || mr1.lock().unwrap().insert(k: "hello2", v: "world2"));

    // can I use Box (smart pointer for heap allocation)?
    let mut map1: HashMap<&str, &str> = HashMap::new();
    map1.insert(k: "hello", v: "world");
    let mr: Arc<Box<HashMap<&str, &str>>> = Arc::new(data: Box::new(map1));
    let mr1: Arc<Box<HashMap<&str, &str>>> = mr.clone();
    thread::spawn(move || (**mr).insert(k: "hello1", v: "world1"));      cannot borrow data in an `Arc` as mutable
    mr1.insert(k: "hello2", v: "world2");      cannot borrow data in an `Arc` as mutable
}
```

## First Principles Thinking



Boiling problems down to their most fundamental truth.

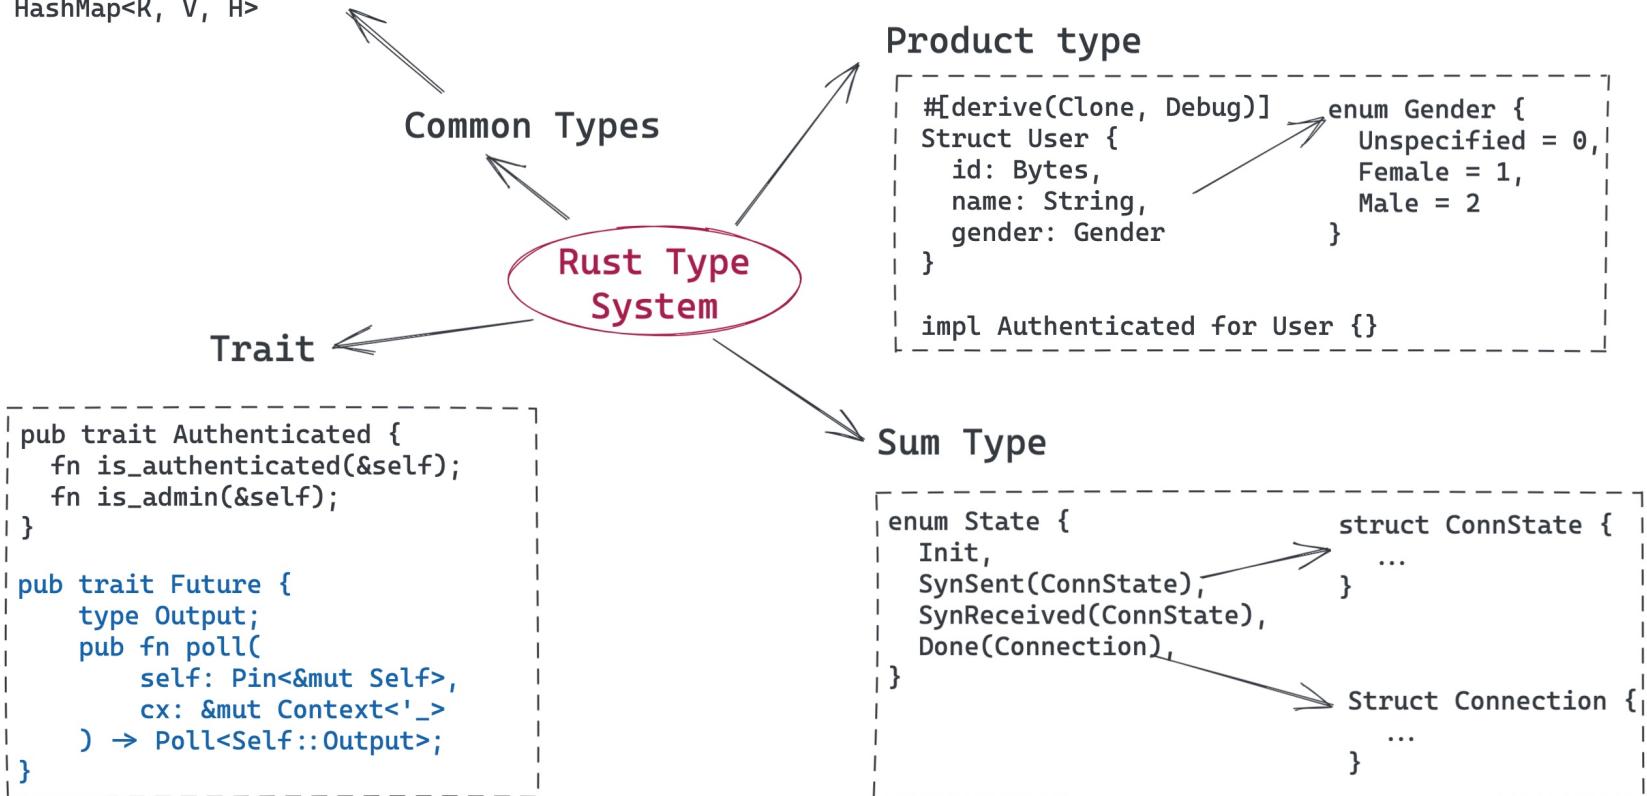
# Recap

- One and only one owner
- Multiple immutable references
- mutable reference is mutual exclusive
- Reference cannot outlive owner
- **use type safety for thread safety**

With these simple rules, Rust achieved safety with  
**zero cost abstraction**

# A glance at Rust Type System

```
Option<T> = T | None  
Result<T, E> = Ok(T) | Err(E)  
Vec<T>  
HashMap<K, V, H>
```



# How's Productivity of Rust?

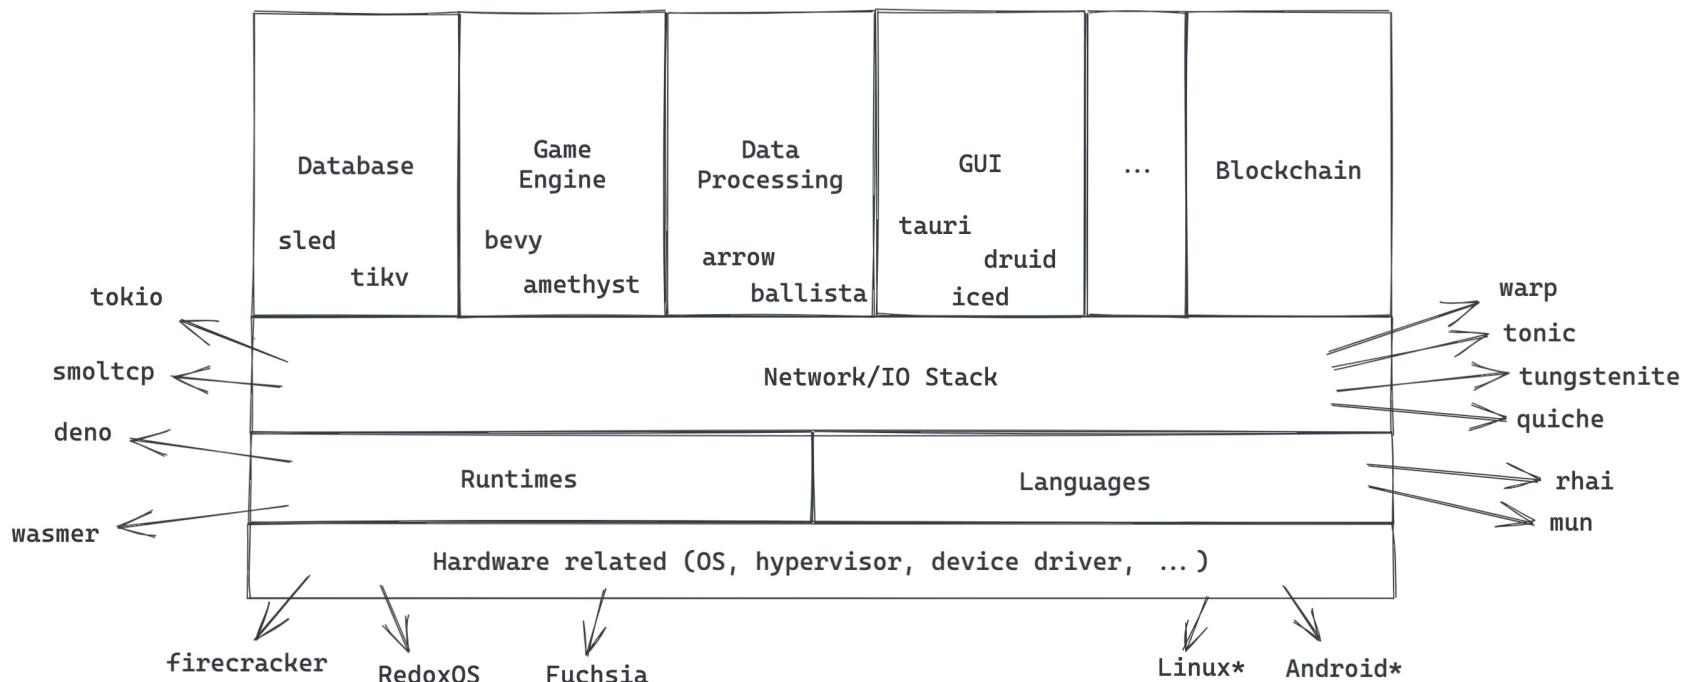
```
1 | # client configuration
2 |
3 | domain = "localhost"
4 |
5 | [cert]
6 | pem = ""-----BEGIN CERTIFICATE-----
7 | MIIBeTCCASugAwIBAgIBKjAFBgMrZXAwNzELMAkGA1UEBgwCVVMxFDASBgNVBAoM
8 | C0RvbWFpbijBjmMuMRIwEAYDVQQDDAlEb21haW4gQ0EwHhcNMjEwMzE0MTg0NTU2
9 | WhcNMzEwMzEyMTg0NTU2WjA3MQswCQYDVQQGDAJVUzEUMBIGA1UECgwLRG9tYWlu
10 | IEluYy4xExAQBgNVBAMMCURvbWFpbjBDQTaqMAUGAytlcAMhAAZhorM9IPsXjBTx
11 | ZxykG15xZrsj3X2XqKjaAVutnf7po1wwjAUBgNVHREEDTALgglsb2Nhbgvc3Qw
12 | HQYDVR00BYEFd+NqChBZD0s5MfMgefHJSIWirthXMBIGA1UDewEB/wQIMAYBaF8C
13 | ARAwDwYDVR0PAQH/BAUDAwcGADAFBgMrZXADQQA9sIlgQcYGaBqTxR1+JadSelMK
14 | Wp35+yhVvuu4PTL18kWdU819w3cVlRe/GHt+jjlbk1i22Tvf05AaNmdxySk0
15 | -----END CERTIFICATE-----"""
16 |
17 | # server configuration
18 |
19 | [identity]
20 | key = ""-----BEGIN PRIVATE KEY-----
21 | MFCAQEwBQYDK2VwBCIEII0kozd0PJsbnfNUS/oqI/Q/enDiLwmdw+JUnTLpR9xs
22 | oSMDIQAtkhJiFdF9SYBIMcLikWPRIgca/Rz9ngIgd6HuG6HI3g==
23 | -----END PRIVATE KEY-----"""
24 |
25 | [identity.cert]
26 | pem = ""-----BEGIN CERTIFICATE-----
27 | MIIBazCCAR2gAwIBAgIBKjAFBgMrZXAwNzELMAkGA1UEBgwCVVMxFDASBgNVBAoM
28 | C0RvbWFpbijBjmMuMRIwEAYDVQQDDAlEb21haW4gQ0EwHhcNMjEwMzE0MTg0NTU2
29 | WhcNMjIwMzE0MTg0NTU2WjA5MQswCQYDVQQGDAJVUzEUMBIGA1UECgwLRG9tYWlu
30 | IEluYy4xFDASBgNVBAMMC0dSUEmgU2VydMVyMcwBQYDK2VwAyEALZISYhXRFuM
31 | SDHC4pFj0SIHGv0c/Z4CIHeh7huhyN6jTDBKMBQGA1UdEQQNMAuCCWxvY2FsaG9z
32 | dDATBgnVHSUEDDAKBgggrBqEFBQcDATAMBgnVHRMEBTADAQEAMA8GA1UdDwEB/wQF
33 | AwMH4AwBQYDK2VwA0EAy7EOIZp73XtcqaSopqDGWU7Umi4DVvIgjmY6qbJZP0sj
34 | ExGdaVq/7M01ZlI+vY7G0NSZWIZUilX0Co0krn0DA==
35 | -----END CERTIFICATE-----"""
36 |
37 |
```

```
9 | ````rust
10 | // you could also build your config with cert and identity separately. See tests.
11 | let config: ServerTlsConfig = toml::from_str(config_file).unwrap();
12 | let acceptor = config.tls_acceptor().unwrap();
13 | let listener = TcpListener::bind(addr).await.unwrap();
14 | tokio::spawn(async move {
15 |     loop {
16 |         let (stream, peer_addr) = listener.accept().await.unwrap();
17 |         let stream = acceptor.accept(stream).await.unwrap();
18 |         info!("server: Accepted client conn with TLS");
19 |
20 |         let fut = async move {
21 |             let (mut reader, mut writer) = split(stream);
22 |             let n = copy(&mut reader, &mut writer).await?;
23 |             writer.flush().await?;
24 |             debug!("Echo: {} - {}", peer_addr, n);
25 |         };
26 |
27 |         tokio::spawn(async move {
28 |             if let Err(err) = fut.await {
29 |                 error!("{}: {:?}", err);
30 |             }
31 |         });
32 |     }
33 | });
34 | ````
```

---

```
36 | Client: You, a month ago • init the project
37 |
38 | ````rust
39 | let msg = b"Hello world\n";
40 | let mut buf = [0; 12];
41 |
42 | // you could also build your config with cert and identity separately. See tests.
43 | let config: ClientTlsConfig = toml::from_str(config_file).unwrap();
44 | let connector = config.tls_connector(Uri::from_static("localhost")).unwrap();
45 |
46 | let stream = TcpStream::connect(addr).await.unwrap();
47 | let mut stream = connector.connect(stream).await.unwrap();
48 | info!("client: TLS conn established");
49 |
50 | stream.write_all(msg).await.unwrap();
51 |
52 | info!("client: send data");
53 |
54 | let (mut reader, _writer) = split(stream);
55 |
56 | reader.read_exact(buf).await.unwrap();
57 |
58 | info!("client: read echoed data");
59 | ````
```

# Things built with Rust



# Should I use Rust?

- Rust is ideal when you need a system that reliable and performant
- Sometimes you don't, sometimes you do, sometimes you need that later
- it's all about tradeoffs

# Rust for our use cases

- parts of the system that are bottlenecks
  - bottleneck on computation
  - bottleneck on memory consumption
  - bottleneck on I/O
- parser/decoder/encoder
- wants to leverage existing C/C++/Rust ecosystem (e.g. you need blake3 for hashing)

# Rust FFI



# Learning rust as a(n)...

- Elixir eng: ownership model, type system, oh no mutation
- Scala eng: ownership model, oh no mutation
- Typescript eng: ownership model, multi-threaded programming
- Swift/Java eng: ownership model
- Python eng: ownership model, type system

**The common  
misunderstandings**

# 1. Rust is super hard to learn...



# Rust is explicit

- Lots of knowledge about computer science is suddenly explicit to you
- If all your pain to learn a lang is 100%:
  - Rust:
    - Compiler help to reduce that to 90%
    - Then you suffer 70% the pains in first 3-6 months
    - Then the rest 20% in 3-5 years
  - Other:
    - You suffer 10-30% in first 3-6 months
    - Then 70%-90% in next 3-5 years

## 2. Unsafe Rust is evil...

Safe Rust	Unsafe Rust
<p>Compiler will make sure dereference is valid</p> 	<p>dereference raw pointer is OK</p> 
<p>Compiler will guarantee it is OK</p> 	<p>deal with FFI</p> <p>impl unsafe trait</p> 
<p>peace of mind</p> <p>The compiler will enforce memory and thread safety</p> 	<p>The dev guarantees memory/thread safety</p> <p>Need extra code review Static analysis</p> 

# References

- The pain of real linear types in Rust
- Substructural type system
- Rust official book
- Rust official site
- Awesome Rust
- Are we web yet?
- Are we async yet?
- Are we gui yet?
- Are we learning yet?
- Are we game yet?
- Are we quantum yet?
- Are we IDE yet?
- Rust is for Professionals

Ownership, borrow check, and  
lifetime

# Ownership/Borrow Rules Review



**Lifetime, not a new idea**

# Lifetime: Stack memory

```

> cat stack.c
#include <stdio.h>
static int VALUE = 42;

void world(char *st, int num) {
    printf("%s(%d)\n", st, num);
}

void hello(int num) {
    char *st = "hello world";
    int v = VALUE+num;
    world(st, v);
}

int main() {
    hello(2);
}

/ttmp
> arm-none-eabi-objdump -d stack.o

stack.o:      file format elf32-littlearm

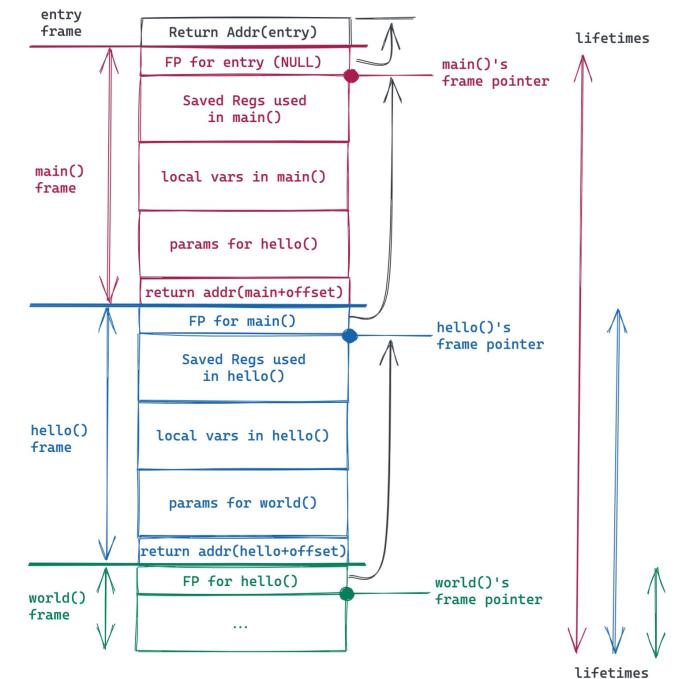
Disassembly of section .text:

00000000 <world>:
  0: e92d4800  push  {fp, lr}
  4: e28db004  add   fp, sp, #4
  8: e24dd008  sub   sp, sp, #8
  c: e50b0008  str   r0, [fp, #-8]
 10: e50b100c  str   r1, [fp, #-12]
 14: e51b200c  ldr   r2, [fp, #-12]
 18: e51b1008  ldr   r1, [fp, #-8]
 1c: e59f0010  ldr   r0, [pc, #16] ; 34 <world+0x34>
 20: ebfffffe  bl    0 <printf>
 24: e1a00000  nop
 28: e24bd004  sub   sp, fp, #4
 2c: e8bd4800  pop   {fp, lr}
 30: e12ffffe  bx    lr
 34: 00000000  .word 0x00000000

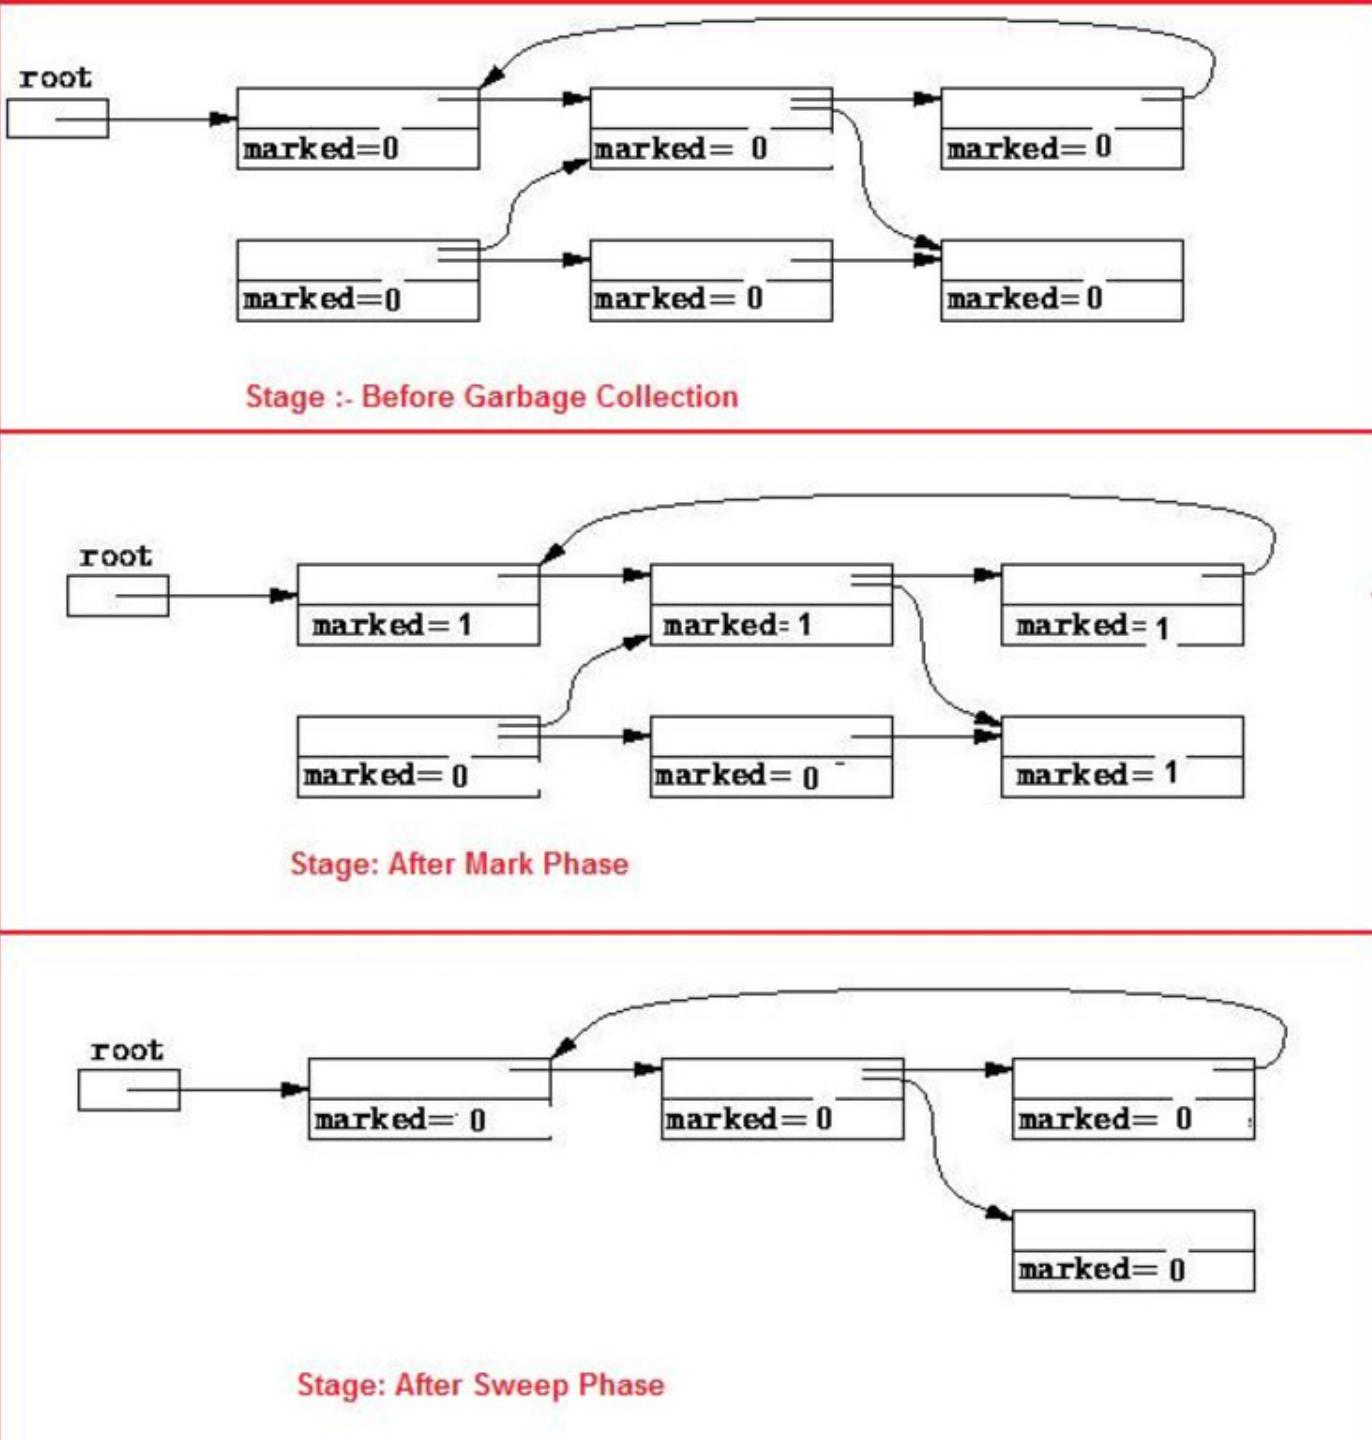
00000038 <hello>:
 38: e92d4800  push  {fp, lr}
 3c: e28db004  add   fp, sp, #4
 40: e24dd010  sub   sp, sp, #16
 44: e50b0010  str   r0, [fp, #-16]
 48: e59f3030  ldr   r3, [pc, #48] ; 80 <hello+0x48>
 4c: e50b3008  str   r3, [fp, #-8]
 50: e59f302c  ldr   r3, [pc, #44] ; 84 <hello+0x4c>
 54: e5933000  ldr   r3, [r3]
 58: e51b2010  ldr   r2, [fp, #-16]
 5c: e0823003  add   r3, r2, r3
 60: e50b300c  str   r3, [fp, #-12]
 64: e51b100c  ldr   r1, [fp, #-12]
 68: e51b0008  ldr   r0, [fp, #-8]
 6c: ebfffffe  bl    0 <world>
 70: e1a00000  nop
 74: e24bd004  sub   sp, fp, #4
 78: e8bd4800  pop   {fp, lr}
 7c: e12ffffe  bx    lr
 80: 00000008  .word 0x00000008
 84: 00000000  .word 0x00000000

00000088 <main>:
 88: e92d4800  push  {fp, lr}
 8c: e28db004  add   fp, sp, #4
 90: e3a00002  mov   r0, #2
 94: ebfffffe  bl    38 <hello>
 98: e3a03000  mov   r3, #0
 9c: e1a00003  mov   r0, r3
  a0: e24bd004  sub   sp, fp, #4
  a4: e8bd4800  pop   {fp, lr}
  a8: e12ffffe  bx    lr

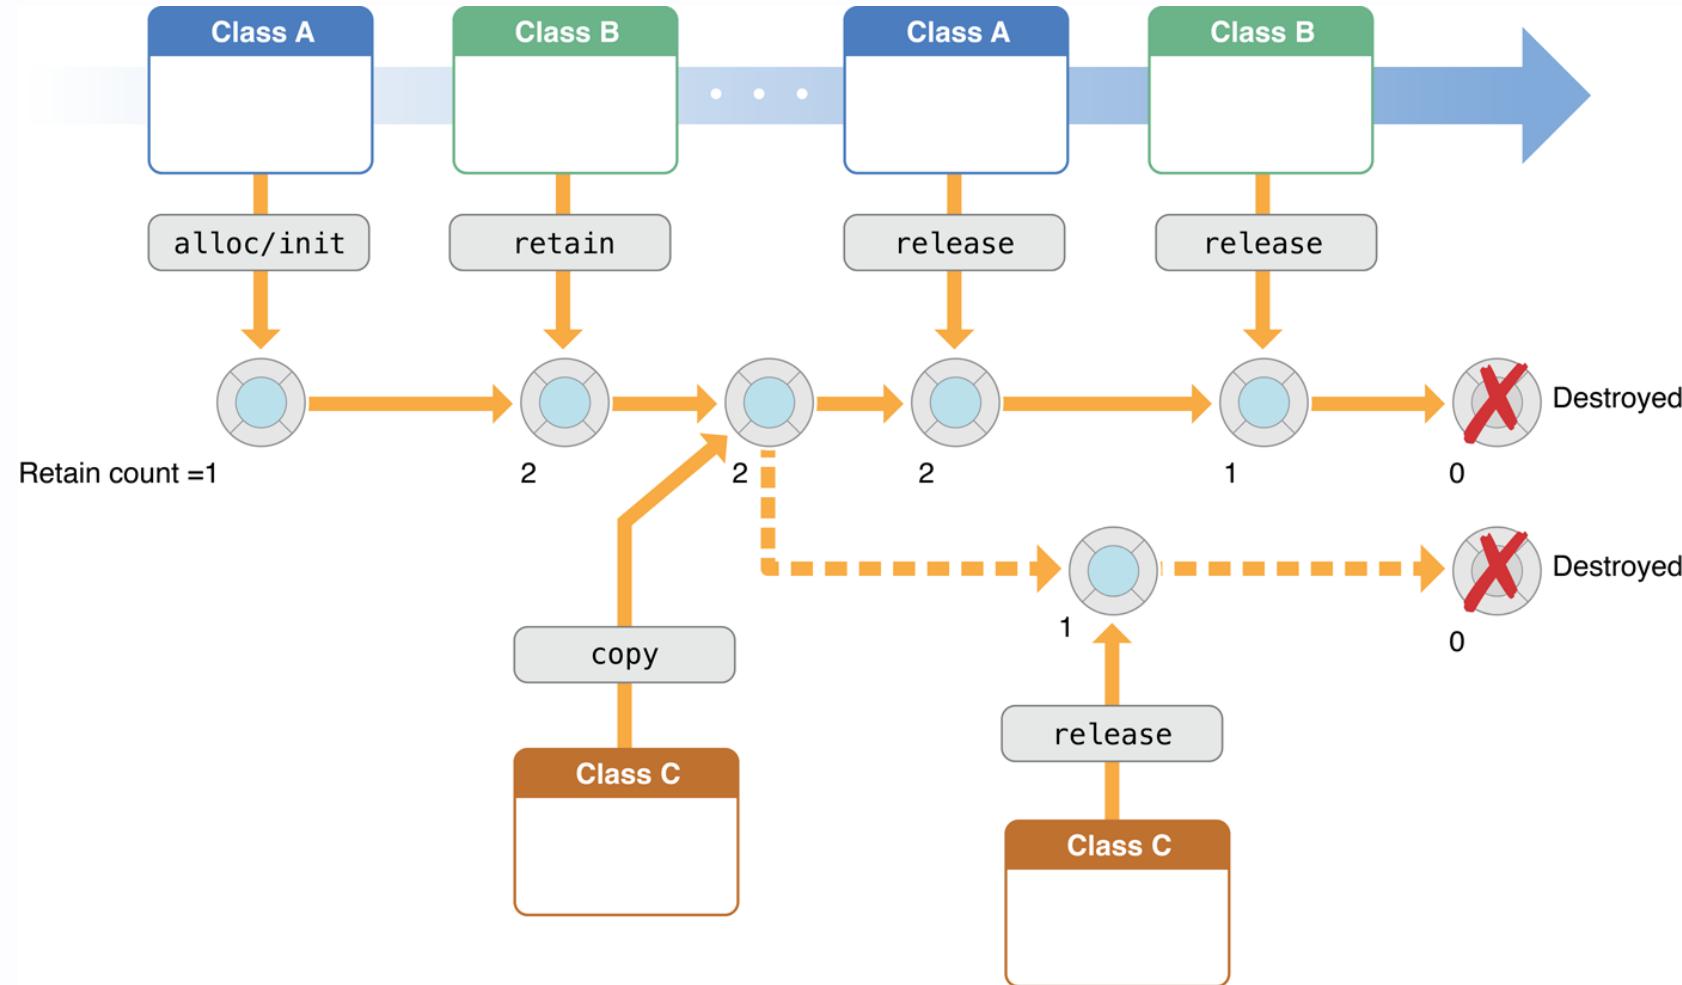
```



# Lifetime: Heap memory (tracing GC)

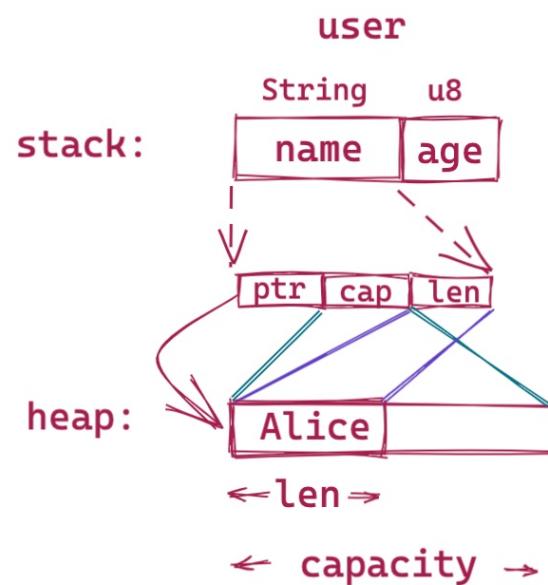


# Lifetime: Heap memory (ARC)

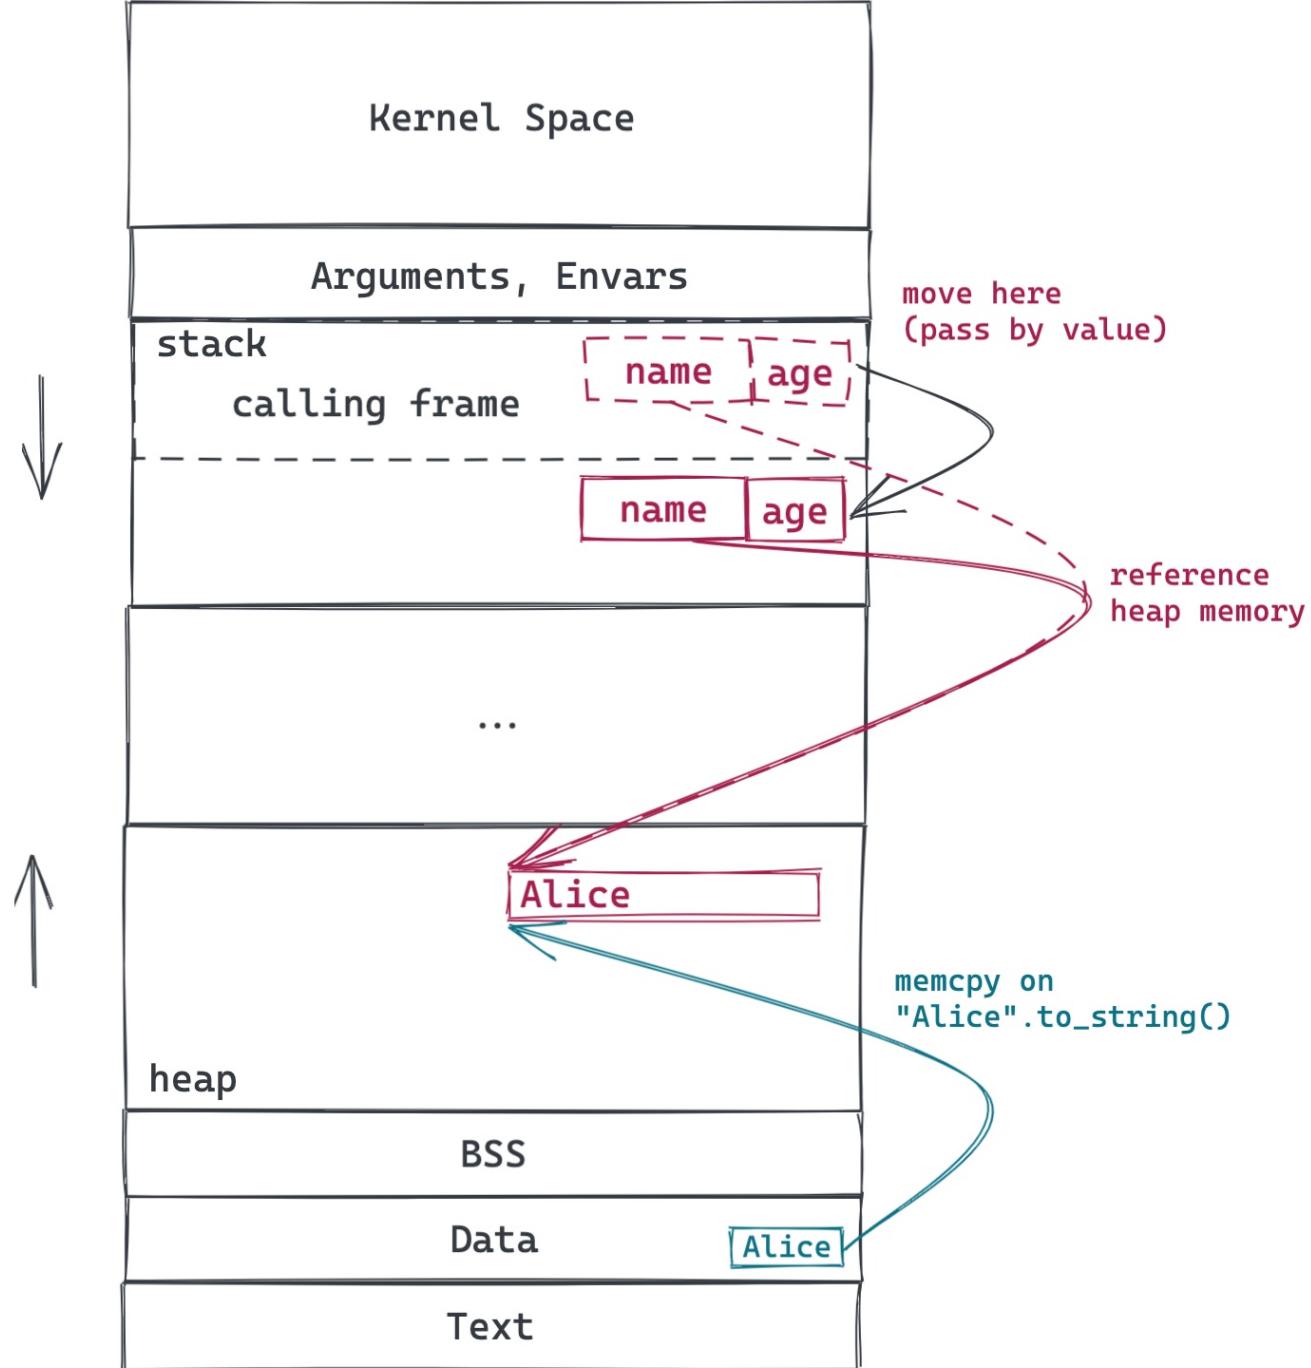


# How Rust handles lifecycle?

# **Move semantics**



```
{
  let user = User {
    name: "Alice".to_string(),
    age: 20,
  };
  let result = insert(user);
  ...
}
```

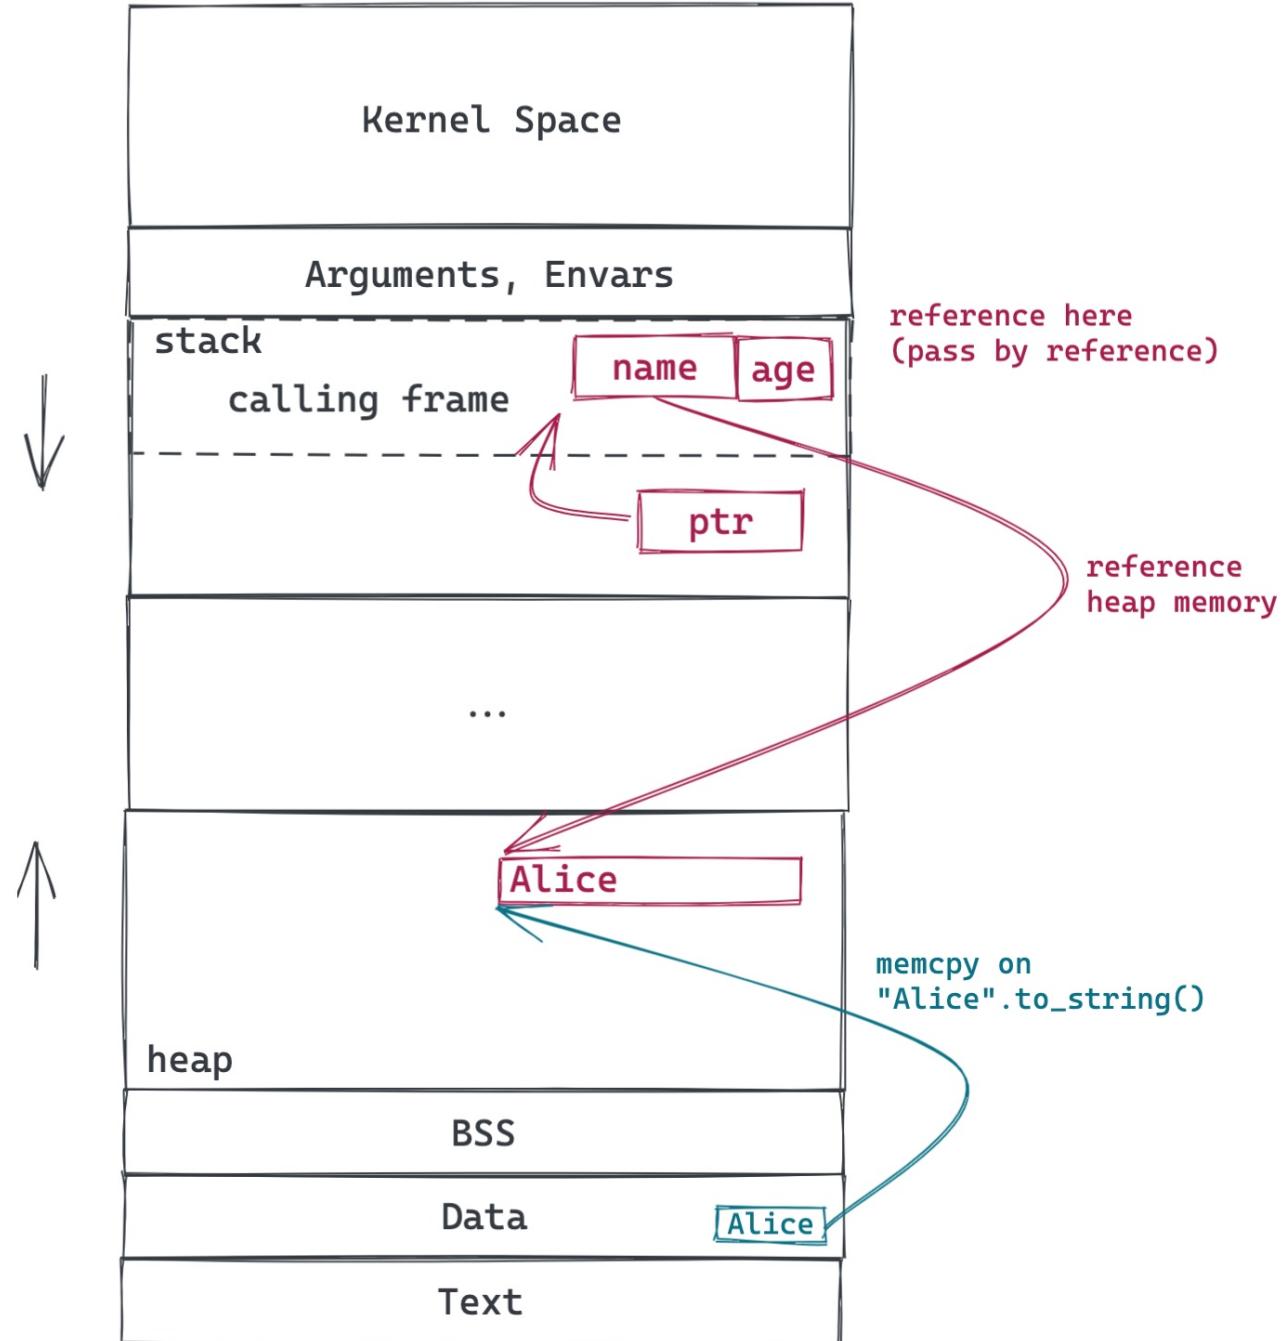


# **Immutable Borrow**

Think about: why does this work in Rust, but not C/C++?



```
{
  let user = User {
    name: "Alice".to_string(),
    age: 20,
  };
  let result = insert(&user);
  ...
}
```



# **Benefit of lifetime-constrained borrow**

- can borrow anything (stack object, heap object)
- safety can be guaranteed at compile-time (no runtime lifetime bookkeeping)
- Rust borrow checker is mostly a lifetime checker

# Lifetime Annotation

- similar as generics, but in lowercase starting with '`'`
- only need to put annotation when there's conflicts

```
// need explicit lifetime
struct User<'a> {
    name: &'a str,
    ...
}
fn process<T, 'a, 'b>(item1: &'a T, item2: &'b T) {}

// &'a User could be written as &User since on confliction
fn lifetime_example(user: &User) { // --- Lifetime 'a
    if user.is_authenticated() { // |--- Lifetime 'b
        let role = user.roles(); // | |
        // | |--- Lifetime 'c
        verify(&role); // | |
        // | |
    } // | ---+
} // ---+}

fn verify(x: &Role) { /*...*/ }
```

# Static Lifetime

- `'static`
- data included in bss / data / text section
  - constants / static variables
  - string literals
  - functions
- if used as trait bound:
  - the type does not contain any non-static references
  - owned data always passes a `'static` lifetime bound, but reference to the owned data does not

# Thread spawn

```
pub fn spawn<F, T>(f: F) -> JoinHandle<T>
where
    F: FnOnce() -> T,
    F: Send + 'static,
    T: Send + 'static,
{
    Builder::new().spawn(f).expect("failed to spawn thread")
}
```

The 'static constraint is a way of saying, roughly, that no borrowed data is permitted in the closure.

# **RAII (Resource Acquisition Is Initialization)**

- initializing the object will also make sure resource is initialized
- releasing the object will also make sure resource is released

# Drop Trait

- memory
- file
- socket
- lock
- any resources

# Mental model

- write the code and defer the complexity about ensuring the code is safe/correct
- confront the most of the safety/correctness problems upfront
- Mutate can only happen when you own the data, or you have a mutable reference
  - either way, the thread is guaranteed to be the only one with access at a time
- Fearless Refactoring
- reinforce properties well-behaved software exhibits
- sometimes too strict: rust isn't dogmatic about enforcing it

# Cost of defects

- Don't introduce defect (this is impossible because humans are fallible).
- Detect and correct defect as soon as the bad key press occurs (within reason: you don't want the programmer to lose too much flow) (milliseconds later).
- At next build / test time (seconds or minutes later).
- When code is shared with others (maybe you push a branch and CI tells you something is wrong) (minutes to days later).
- During code review (minutes to days later).
- When code is integrated (e.g. merged) (minutes to days later).
- When code is deployed (minutes to days or even months later).
- When a bug is reported long after the code has been deployed (weeks to years later).

# Ownership and Borrow rules

- Use after free: no more (reference can't point to dropped value)
- Buffer underruns, overflows, illegal memory access: no more (reference must be valid and point to owned value)
- memory level data races: no more (single writer or multiple readers)

# References

- [Mark-And-Sweep \(Garbage Collection Algorithm\)](#)
- [Tracing garbage collection](#)
- [Swift: Avoiding Memory Leaks by Examples](#)
- [Reference counting](#)
- [Fearless concurrency with Rust](#)
- [Rust means never having to close a socket](#)

# Typesystem and data structures

- The type system is relatively strong and prevents many classes of bugs.
- The borrow checker and the rules it enforces prevent safety issues at compile time. Some of these violations can be detected by other languages' compilers. However, in many cases sufficient auditing (like {address, memory, thread} sanitizers) is run much less frequently, often only in CI tests, which can be hours or days later.
- Invariants can be encoded and enforced in the type system through features like enums being algebraic data types.
- Variables are immutable by default and must be explicitly annotated as mutable. This forces you to think about where and how data mutation occurs, enabling you to spot issues sooner.
- Option<T> significantly curtails the billion dollar mistake.
- Result<T, E> forces you to reckon about handling errors.

# Concurrency - primitives

# Concurrency - `async/await`

# Networking and security

# FFI with C/Elixir/Swift/Java

# WASM/WASI

# Rust for real-world problems

May the **Rust** be with you